

# QBR narratives, grounded in data

Prepared for Sterling FinTech Ltd · FinTech

Turn verified usage metrics into review-ready QBR narratives, with a CSM accountable for every customer-facing draft.

## 01 / Outcomes

- Reduce CSM preparation time while preserving account-specific judgement.
- Ensure every numerical claim reconciles to a verified metric snapshot.
- Prove the value and controls needed for a wider rollout.

## 02 / Scope boundaries

- In: One agreed usage-data export and one QBR narrative template.
- In: Drafts with metric references, period labels and missing-data flags.
- In: SSO for pilot users and an approved deployment region.
- In: A pilot with 5 CSMs and up to 10 fictional or approved test accounts.
- Out: Autonomous customer delivery or CRM write-back.
- Out: Predictive revenue claims or invented benchmarks.
- Out: Replacing dashboards, rebuilding the data warehouse or all QBR templates.
- Out: Organisation-wide rollout, multilingual output and ongoing managed service.

## 03 / Delivery approach

- Discover (1–2 weeks): Confirm authoritative metrics, account mapping, residency and the CSM baseline. Activities: Interview the sponsor and nominated pilot users; Profile approved sample data and identify access gaps; Agree baseline measures, acceptance thresholds and a go/no-go gate. Exit: Sponsor signs off the pilot boundary, data access and evaluation plan.
- Design (1–2 weeks): Turn agreed outcomes and controls into a buildable pilot design. Activities: Map the bounded workflow and reviewer responsibilities; Define interfaces, failure handling and evidence requirements; Review the design with the operational and compliance owners. Exit: Design and required control owners approve the implementation plan.
- Build (2–4 weeks): Build metric-grounded narratives, numerical checks, SSO and approval workflow. Activities: Implement only the agreed data inputs and workflow; Add traceability, access controls and explicit failure states; Evaluate representative and adversarial cases against the reference set. Exit: The test cohort meets agreed quality thresholds; critical issues are resolved.
- Pilot (2–3 weeks): Evaluate usefulness and reliability with the agreed user cohort. Activities: Onboard pilot users and train reviewers; Run a bounded trial against the baseline; Capture errors, review effort and user feedback. Exit: Sponsor reviews measured results and decides to stop, iterate or expand.
- Handover (1 week): Give the owner a controlled, maintainable next step. Activities: Document the workflow, limits and operational runbook; Transfer configuration and review responsibilities; Present the pilot findings and a separately scoped rollout option. Exit: Named owner accepts the handover and the next-step decision is recorded.

## 04 / Planning estimate

- Medium engagement · 12–23 person-weeks. Planning estimate, not a quote; elapsed time depends on staffing and dependencies.
- One usage export and identity provider; no CRM write-back.
- Metric definitions, periods and account matching need validation.
- SSO configuration and role access for the pilot cohort.
- Approved processing region and retention settings.
- CSMs approve each narrative before customer use.

## 05 / Assumptions to confirm

- Low: Usage exports contain stable account IDs and reporting periods.
- Medium: An identity owner can support the pilot SSO integration.
- High: CSMs retain approval of every customer-facing narrative.

## 06 / Decisions needed

- Commercial: What pilot budget and preparation-time saving would justify expansion? Sets an investment gate before committing to rollout.
- Technical: Which source is authoritative when the export and CRM disagree? Prevents contradictory numerical claims and defines reconciliation rules.
- Operational: Who approves the narrative template and reviews the first 10 drafts? Makes accountability and acceptance criteria explicit.
- Compliance: Which hosting region, retention period and model-processing terms are required? Data residency is a design constraint, not a post-build checkbox.

## 07 / Risks & controls

- The model invents or misstates a number. Pass only verified metrics, validate numerical claims and block unsupported drafts.
- Inconsistent reporting periods distort the story. Freeze a metric dictionary and flag incomparable or missing periods.
- Residency requirements constrain provider choice. Confirm deployment and processing terms before selecting a live model.

## 08 / Measuring success

- Preparation time: Proposed: 50% reduction against a two-week baseline.
- Numerical accuracy: Proposed: 100% of numerical claims reconcile to the supplied snapshot.
- CSM usefulness: Proposed: 80% of drafts rated useful with minor edits.